

Causal Rejected-Side Masking for Reasoning Preference Optimization

Anonymous authors
Paper under double-blind review

Abstract

Direct Preference Optimization (DPO) penalizes rejected responses as whole trajectories. This is a poor inductive bias for multi-step reasoning: a rejected solution can contain a correct prefix, a localized first error, and later steps that merely inherit the wrong state. We study causal rejected-side masking, a family of DPO-style objectives that changes only how the rejected response is aggregated. The central operation is prefix protection: steps before the first error are masked out of the rejected-side penalty. Causal-Masked DPO (CM-DPO) further adds a decayed penalty on post-error suffix steps. Controlled arithmetic experiments show that vanilla DPO strongly degrades reasoning accuracy, while causal masking improves retention and held-out performance. Across three harder-arithmetic seeds, CM-DPO is best on average, but only narrowly over a first-error-only baseline. GSM8K-style pilots show a mixed result: prefix protection consistently helps, but suffix decay has not yet produced a reliable additional gain. We therefore present CM-DPO as a causal masking framework, with prefix protection as the robust empirical finding and suffix decay as a promising but not yet settled design choice.

1 Introduction

Preference optimization is widely used to adapt language models to desired behavior. Direct Preference Optimization (DPO) (Rafailov et al., 2023) is particularly attractive because it trains directly on preferred and rejected responses without fitting a separate reward model. However, the standard DPO objective treats a rejected response as uniformly negative evidence. This is a natural simplification for short answers, but a questionable one for long reasoning traces.

The problem is causal credit assignment. A rejected mathematical solution can correctly understand the question, set up the right equation, and then make one local arithmetic mistake. Penalizing the entire rejected trajectory suppresses both the actual error and the correct reasoning prefix. This creates an over-penalty failure mode: the model is trained away from useful intermediate reasoning because that reasoning happened to appear inside a rejected final answer.

We study a simple alternative: causal rejected-side masking. Given a rejected response segmented into steps and an estimate of the first erroneous step, we change only the rejected-side aggregation in DPO. Steps before the first error are protected from negative pressure. The first error receives full penalty. Later steps can either be fully penalized, ignored, or decayed. Causal-Masked DPO (CM-DPO) is one point in this family, using prefix protection plus exponential suffix decay. The framework keeps the same preference-pair interface as DPO and requires only step boundaries and a first-error estimate.

Our empirical results suggest a more nuanced story than the strongest version of the initial hypothesis. On controlled harder arithmetic data, causal masking clearly avoids the severe degradation caused by vanilla DPO. In the strongest multi-seed setting, CM-DPO achieves the best mean held-out accuracy. At the same time, its margin over first-error-only DPO

is small. On GSM8K-style pilots, prefix protection remains useful, but suffix decay has not yet shown a reliable benefit.

We therefore frame the contribution conservatively. The robust empirical claim is not that one fixed suffix-decay schedule is universally best. Rather, rejected reasoning trajectories contain causal structure, and DPO benefits from respecting at least the most important part of that structure: correct prefixes should not be penalized as if they were errors.

Our contributions are:

- We formulate trajectory-level DPO over-penalty as a causal credit-assignment problem in multi-step reasoning.
- We introduce a rejected-side masking family that includes vanilla DPO, prefix-masked DPO, first-error-only DPO, and CM-DPO as different weight choices under one objective.
- We implement token-level masked DPO with LoRA training and first-error localization, enabling direct comparisons under the same model, data, and optimization setup.
- We report controlled arithmetic and GSM8K-style pilots showing that prefix protection is robustly useful, while suffix decay remains task- and data-quality-sensitive.

2 Related Work

Preference optimization. Reinforcement learning from human feedback aligns language models by learning a reward model and optimizing against it (Ouyang et al., 2022). DPO (Rafailov et al., 2023) removes the explicit reward-model stage and directly optimizes preference pairs through policy-reference likelihood ratios. Our work keeps DPO’s trajectory-level data interface, but modifies how rejected trajectories contribute to the objective. This makes the proposal orthogonal to the source of the preference pair: it can be applied to human preferences, verifier-generated preferences, or synthetic reasoning preferences when a first-error signal is available.

Reasoning supervision. Benchmarks such as GSM8K (Cobbe et al., 2021) and MATH (Hendrycks et al., 2021) emphasize multi-step mathematical reasoning. Process supervision and step-level verification aim to supervise intermediate reasoning rather than only final answers (Lightman et al., 2024). Causal rejected-side masking is complementary: it can consume first-error signals from verifiers or judges, but does not require transforming the data into explicit step-level preference pairs.

Credit assignment in reasoning traces. Outcome-level supervision gives a single label to an entire solution. This is convenient, but it conflates the step that caused failure with earlier steps that may be correct and later steps that may be consequences of the first error. Process supervision addresses this by labeling intermediate states. Our setting lies between these extremes: we assume a first-error localization signal, then use it to reweight the rejected side of a standard preference objective. The goal is not to replace process supervision, but to use limited process information inside a DPO-compatible loss.

Masked and weighted objectives. The method can be viewed as a token-weighted variant of DPO in which weights are induced by step-level causal structure. The preferred response remains fully weighted, while the rejected response receives structured weights. This differs from generic loss reweighting because the mask is asymmetric and causal: tokens before the first error in the rejected response are not treated as negative evidence.

Parameter-efficient adaptation. Our experiments use LoRA (Hu et al., 2022) to make repeated preference-optimization comparisons feasible. The masking objective is independent of LoRA and can also be used with full fine-tuning.

Table 1: Rejected-side causal masking family. The preferred response is always fully weighted.

Variant	Prefix $k < m$	First error $k = m$	Suffix $k > m$
Vanilla DPO	1	1	1
Prefix-masked DPO	0	1	1
First-error-only DPO	0	1	0
CM-DPO	0	1	γ^{k-m}

3 Method

3.1 DPO and rejected-side over-penalty

For a prompt x , DPO uses a preferred response y_w and rejected response y_l . Its implicit reward is

$$r_\theta(y | x) = \log \pi_\theta(y | x) - \log \pi_{\text{ref}}(y | x), \quad (1)$$

and the loss is

$$\mathcal{L}_{\text{DPO}}(\theta) = -\mathbb{E}[\log \sigma(\beta(r_\theta(y_w | x) - r_\theta(y_l | x)))]. \quad (2)$$

When y_l is a multi-step reasoning trace, this objective penalizes all rejected tokens even if some earlier steps are correct. The failure is specific to the rejected side: the preferred response should remain positive evidence, but a rejected response may contain both positive and negative reasoning evidence.

3.2 Causal step weights

Let the rejected response be segmented into steps

$$y_l = (s_1, \dots, s_K), \quad (3)$$

and let m be the estimated first-error index. A causal rejected-side mask assigns each rejected step a nonnegative weight w_k . Steps with $k < m$ occurred before the localized error and should not receive negative pressure from the rejected label. The preferred response remains fully weighted. The weighted rejected reward is

$$r_\theta^w(y_l | x) = \sum_{k=1}^K w_k [\log \pi_\theta(s_k | x, s_{<k}) - \log \pi_{\text{ref}}(s_k | x, s_{<k})]. \quad (4)$$

The masked DPO loss is

$$\mathcal{L}_{w\text{-DPO}}(\theta) = -\mathbb{E}[\log \sigma(\beta(r_\theta(y_w | x) - r_\theta^w(y_l | x)))]. \quad (5)$$

3.3 Variants

Table 1 shows the four weight patterns used in the experiments. These variants isolate two design choices: whether to protect the pre-error prefix, and how much penalty to assign after the first error.

CM-DPO therefore adds suffix decay to the robust prefix-protection idea. This distinction matters empirically: the prefix mask is consistently useful, while the extra decay design is not always better than simpler alternatives.

3.4 Implementation

Step-level weights are expanded to token-level response masks by aligning segmented step spans to response tokens. The DPO trainer computes policy and reference log-likelihoods with these masks applied to the rejected response. The reference model is frozen, and all reported model updates use LoRA adapters. If a rejected response cannot be localized or if the first error is ambiguous, the example is filtered rather than assigned a high-confidence causal mask.

4 Experiments

4.1 Controlled harder arithmetic

The main mechanism experiments use synthetic harder arithmetic preference pairs. Each example contains a prompt, a correct step-by-step solution, and a rejected solution with a localized arithmetic or state-transition error. This setting gives high-quality first-error labels and directly tests whether rejected-side masking protects correct prefixes.

We train Qwen2.5-0.5B-Instruct with LoRA adapters. The key scaled setting uses 2,000 preference pairs, $\gamma = 0.25$, one epoch, $\beta = 0.1$, maximum length 768, batch size 1, and gradient accumulation 8. We evaluate on 500 held-out harder arithmetic problems.

The controlled setting is intentionally narrow. It is designed to test the mechanism under clean localization rather than to claim broad mathematical ability. This makes failures informative: if vanilla DPO harms accuracy even when the preference data are clean, then the problem is not only noisy labels but also the way rejected reasoning is aggregated.

4.2 GSM8K-style pilots

We also run GSM8K-style pilots to test transfer beyond controlled templates. We generate preference pairs from GSM8K prompts using the base model, select sampled correct responses when available, and otherwise fall back to the gold answer. Rejected solutions are localized using an API judge that returns the first erroneous step and confidence. We filter examples by confidence, non-final first-error position, and truncation heuristics.

The GSM8K pilots are intentionally treated as diagnostic rather than final benchmark evidence. They reveal whether the same masking idea survives noisier generated reasoning, but they also expose data-quality issues such as gold-fallback chosen responses and style mismatch.

4.3 Evaluation questions

We evaluate the masking family through three questions. First, does vanilla DPO degrade reasoning behavior when rejected traces contain correct prefixes? Second, does masking the pre-error prefix recover performance? Third, does assigning a decayed penalty to post-error suffix steps improve over simpler first-error-only or prefix-masked alternatives? The first two questions are the core causal-masking claims; the third is the specific additional hypothesis behind CM-DPO.

5 Results

5.1 Harder arithmetic supports causal masking

Table 2 shows the 500-example harder arithmetic evaluation after training on 2,000 preference pairs over three seeds. Vanilla DPO is consistently harmful. Prefix-masked and first-error-only DPO are much stronger. CM-DPO is best on average, but its margin over first-error-only DPO is small.

This result supports the central causal-masking claim: uniformly penalizing rejected reasoning is unsafe, and protecting pre-error reasoning substantially improves behavior. It does not, by itself, prove that suffix decay is much better than simply training on the first erroneous step.

5.2 Ablation: prefix protection versus suffix decay

The four variants decompose the effect of causal masking. Vanilla DPO has no prefix protection and drops from a base accuracy of 0.410 to 0.253. Prefix-masked DPO changes only the pre-error prefix weights and recovers much of the lost performance, reaching 0.365.

Table 2: Harder arithmetic eval500 after harder2000 training. Mean and standard deviation are over three seeds.

Model	Accuracy	Std.
Base	0.410	–
Vanilla DPO	0.253	0.054
Prefix-masked DPO	0.365	0.012
First-error-only DPO	0.427	0.016
CM-DPO	0.433	0.011

Table 3: GSM8K-style eval100. Entries are correct out of 100.

Training data	Base	Vanilla	Prefix	First-error	CM-DPO
Cheap localizer, 49 rows	16	19	22	21	20
API judge, 49 rows	16	18	25	23	20
API judge HQ, 109 rows	16	25	29	28	28
More-candidates HQ, 87 rows	16	22	25	28	27

First-error-only DPO removes suffix pressure entirely and reaches 0.427. CM-DPO keeps a decayed suffix penalty and reaches 0.433.

This pattern supports a clear primary conclusion and a cautious secondary one. The primary conclusion is that prefix protection is responsible for the largest improvement over vanilla DPO. The secondary conclusion is that suffix decay can help in the controlled setting, but the observed gain over first-error-only DPO is small relative to the current scale of the experiments.

5.3 GSM8K pilots are mixed

Table 3 summarizes the GSM8K-style pilots on a 100-example evaluation set. As localization and filtering improve, all trained variants can beat the base model. However, CM-DPO does not dominate. Prefix-masked DPO is best in the 109-example high-quality setting, while first-error-only is best in the more-candidates setting.

5.4 Localization and data quality

Table 4 shows why the GSM8K pilots should be interpreted cautiously. API judge localization greatly improves the fraction of examples with non-final first errors, but strict high-quality filtering still leaves limited data. Increasing candidates raises the sampled-correct chosen rate in raw data, yet does not improve downstream accuracy after filtering.

The robust conclusion is that prefix protection matters. The suffix-decay component of CM-DPO remains plausible on controlled arithmetic, but the GSM8K pilots do not yet establish it as a reliable additional improvement.

6 Discussion

The experiments suggest separating the method into two claims. The first claim is prefix protection: rejected reasoning trajectories should not be uniformly penalized when their prefixes are correct. This claim is supported across controlled arithmetic and GSM8K-style pilots. The second claim is suffix decay: after the first error, later steps should receive a decayed but nonzero penalty. This claim is supported by some controlled likelihood evidence, but not yet by the GSM8K pilots.

Table 4: GSM8K-style localization and filtering diagnostics.

Dataset	Rows	Non-final first error	Sampled chosen
Cheap localizer	49	1 / 49	26 / 49
API judge	49	26 / 49	26 / 49
API judge raw	199	121 / 199	77 / 199
API judge HQ	109	109 / 109	34 / 109
More-candidates raw	200	101 / 200	101 / 200
More-candidates HQ	87	87 / 87	35 / 87

This distinction changes the paper’s center of gravity. Instead of presenting CM-DPO as a fixed objective whose suffix schedule is always best, we present causal rejected-side masking as a framework. CM-DPO is one instance of the framework. Prefix-masked and first-error-only objectives are not merely weak baselines; they are important ablations that reveal which part of the causal mask is doing work.

For future experiments, the most useful direction is not simply more GSM8K data. The data must contain style-matched chosen and rejected responses, high-confidence first-error labels, and enough post-error suffix content for suffix decay to matter. Without those properties, GSM8K training mostly tests task adaptation and prefix protection, not the full CM-DPO hypothesis.

This framing also affects how negative results should be interpreted. If CM-DPO does not beat first-error-only DPO on a noisy dataset, that does not necessarily refute causal masking. It may indicate that the available suffix tokens are low-value training signal, that localization is too noisy, or that the model benefits most from correcting the first error directly. The framework makes these possibilities experimentally separable.

7 Limitations

The current evidence is preliminary. The strongest positive results come from controlled harder arithmetic, where first-error labels are cleaner than in open-ended model generations. This makes the mechanism easy to test but does not prove broad reasoning generalization.

The GSM8K pilots are intentionally reported as mixed evidence. They show that the pipeline runs end-to-end and that prefix protection helps, but they do not show a reliable suffix-decay benefit. Many high-quality filtered examples still use gold-fallback chosen responses, creating a style mismatch with model-generated rejected responses.

The current experiments also use Qwen2.5-0.5B-Instruct with LoRA. Larger models, stronger sampled solution generators, and better first-error judges may change the balance between prefix-masked, first-error-only, and CM-DPO objectives.

Finally, the paper does not yet include standard full-scale GSM8K or MATH evaluations. Those should be added only after the preference data construction yields enough style-matched sampled-correct chosen responses. The current draft should therefore be read as a mechanism paper in progress, not as a complete benchmark claim.

8 Conclusion

We studied causal rejected-side masking for reasoning preference optimization. The central empirical finding is that DPO should not treat partially correct rejected reasoning as uniformly negative: protecting correct prefixes consistently improves behavior relative to vanilla DPO. CM-DPO extends this idea with suffix decay and performs best on average in controlled harder arithmetic, but GSM8K-style pilots show that suffix decay is not yet a robust advantage over simpler masks. Future work should focus on cleaner style-matched preference data and stronger first-error localization before claiming broad GSM8K gains.

References

- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. Training verifiers to solve math word problems. In arXiv preprint arXiv:2110.14168, 2021.
- Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song, and Jacob Steinhardt. Measuring mathematical problem solving with the MATH dataset. In Thirty-fifth Conference on Neural Information Processing Systems Datasets and Benchmarks Track, 2021.
- Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. In International Conference on Learning Representations, 2022.
- Hunter Lightman, Vineet Kosaraju, Yuri Burda, Harrison Edwards, Bowen Baker, Teddy Lee, Jan Leike, John Schulman, Ilya Sutskever, and Karl Cobbe. Let’s verify step by step. In International Conference on Learning Representations, 2024.
- Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll L. Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, John Schulman, Jacob Hilton, Fraser Kelton, Luke Miller, Maddie Simens, Amanda Askell, Peter Welinder, Paul Christiano, Jan Leike, and Ryan Lowe. Training language models to follow instructions with human feedback. In Advances in Neural Information Processing Systems, 2022.
- Rafael Rafailov, Archit Sharma, Eric Mitchell, Christopher D. Manning, Stefano Ermon, and Chelsea Finn. Direct preference optimization: Your language model is secretly a reward model. In Advances in Neural Information Processing Systems, 2023.